When Agents Disagree With Themselves: Measuring Behavioral Consistency in LLM-Based Agents

Aman Mehta [1]

Abstract

Run the same LLM agent on the same task twice: do you get the same behavior? We find the answer is often no. In a study of 10,000 agent runs across five models and four providers on 200 HotpotQA questions, ReAct-style agents produce 2.3–4.2 distinct action sequences per 10 runs on average. This variance is associated with failure: consistent tasks ($\leq$2 unique paths) achieve 80–87% accuracy while inconsistent tasks ($\geq$4 paths) achieve 43–65%, a gap that persists after controlling for task difficulty (partial $r = -0.17$ to -0.50, $p < 0.02$). Divergence concentrates at step 2 (52.5% of Llama tasks), and answer entropy achieves AUROC 0.67–0.69 as a lightweight failure detector. Exploiting this signal, selective prediction—answering only when $k = 3$ runs agree—achieves 87–88% accuracy at 54–62% coverage, a 6–14pp gain over single-run baselines. A cross-benchmark pilot on SWE-bench (150 runs) preserves the consistency hierarchy, and bootstrap analysis shows single-run evaluations produce incorrect model rankings 70.1% of the time. Our results suggest that multi-run evaluation and consistency monitoring should become standard practice for agent deployments.

1. Introduction

Large language model (LLM) based agents that use tools and multi-step reasoning are increasingly deployed for complex tasks (Yao et al., 2022; Schick et al., 2023). As these systems move from prototypes to production, understanding their reliability becomes critical.

A fundamental but understudied question is: *how consistent are LLM agents in their behavior?* Given identical inputs, will an agent follow the same reasoning path and arrive at the same answer? This matters because inconsistency complicates debugging, may signal uncertainty exploitable for error detection, and could inform architecture decisions.

We present a systematic empirical study of behavioral consistency in ReAct-style agents (Yao et al., 2022). Our key contributions:

1. **Quantified behavioral variance**: 10,000 runs (200 tasks $\times$ 10 runs $\times$ 5 models) reveal 2.3–4.2 unique action sequences per task on average.

2. **Consistency–correctness association**: Consistent tasks achieve 80–87% accuracy vs. 43–65% for inconsistent tasks, with effect sizes $r = -0.35$ to -0.58 ($p < 0.001$) that hold after controlling for task difficulty (partial $r = -0.17$ to -0.50).

3. **Failure detection**: Answer entropy achieves AUROC 0.67–0.69 as a runtime failure signal, with precision up to 67% vs. a 17–24% baseline.

4. **Actionable intervention**: Selective prediction via unanimous $k = 3$ agreement achieves 87–88% accuracy (6–14pp above baselines) at 54–62% coverage, instantiating the selective classification framework for agents.

5. **Early divergence and path length**: 20.0% of tasks diverge by step 2 (52.5% for Llama); path length correlates with both consistency and correctness ($\rho = 0.43$–0.74).

6. **Ranking instability**: Single-run evaluations produce incorrect model rankings 70.1% [69.2–71.0%] of the time, undermining standard benchmark methodology.

2. Related Work

LLM-Based Agents. ReAct (Yao et al., 2022) introduced interleaving reasoning with actions. Subsequent work extended this to web browsing (Zhou et al., 2023), code execution (Jimenez et al., 2023b), and multi-agent systems (Wu et al., 2023). We study reliability under repeated execution.

[1]Snowflake AI Research. Correspondence to: Aman Mehta <aman.mehta@snowflake.com>.

Preprint. March 16, 2026.

Agent Consistency. τ-bench (Yao et al., 2024) showed that GPT-4o's pass rate drops from 60% (pass^1) to 25% (pass^8). While τ-bench quantifies *whether* agents are inconsistent, we investigate *where* and *why*: tracing divergence to step 2, correlating consistency with correctness, and identifying path length as a signal. Stroebl et al. (2024) argue that benchmarks overestimate agent capabilities; our findings support this.

Uncertainty and Selective Prediction. Our work connects to three threads. First, *ensemble disagreement*: deep ensembles use prediction variance as an uncertainty estimate (Lakshminarayanan et al., 2017); behavioral consistency across agent runs is an analogous signal, where each run acts as an implicit ensemble member. Second, *self-consistency*: Wang et al. (2023) show that majority voting over sampled reasoning chains improves accuracy; we extend this from single-turn CoT to multi-step agentic settings and additionally study consistency as a *diagnostic* signal, not just an ensembling strategy (Section 4.4). Third, *selective prediction*: Geifman & El-Yaniv (2017) and Kamath et al. (2020) show that abstaining on uncertain inputs improves precision; our consistency-filtered results instantiate this for agents. Unlike verbalized confidence (Xiong et al., 2024; Tian et al., 2023), behavioral consistency requires no model modification—it emerges naturally from repeated execution. Prior work on LLM calibration (Kadavath et al., 2022) and semantic uncertainty (Kuhn et al., 2023) focuses on single-turn outputs; we extend to multi-step agentic behavior.

3. Methodology

3.1. Agent Architecture

We implement a ReAct-style agent with three tools: **Search(query)** returns document titles via keyword matching, **Retrieve(title)** returns full text, and **Finish(answer)** terminates with a final answer. The agent follows the standard think–act–observe loop.

3.2. Experimental Setup

Dataset. We use 200 questions from HotpotQA (Yang et al., 2018) validation (distractor setting), requiring multi-hop reasoning over 10 paragraphs (2 gold, 8 distractors).

Models. Five models across four providers: Llama 3.1 70B (Meta), GPT-4o and GPT-5 (OpenAI), Claude Sonnet 4.5 (Anthropic), and Gemini 3 Pro (Google).

Runs. For each question–model pair, 10 independent runs at temperature 0.7, yielding 10,000 total runs. Temperature ablation: 20 questions $\times$ 5 runs $\times$ 4 stochastic models $\times$ 3 temperatures.

Statistical note. All tests are two-sided; raw p-values without multiple-comparison correction given the exploratory nature.

3.3. Metrics

Answer consistency: fraction of runs producing the most common answer. **Action sequence diversity**: number of unique action sequences across N runs. **Step variance ratio**: $(\max - \min)/\text{mean}$ of step counts. **First divergence point**: earliest step where runs take different actions. **Correctness**: fuzzy string matching (answer contains gold or vice versa, case-insensitive); results under exact match and token F1 in Appendix F.

4. Results

4.1. Overall Model Comparison

Table 1 summarizes 10,000 runs. Claude Sonnet 4.5 achieves both the highest accuracy and consistency (2.3 unique sequences). Llama 3.1 70B shows the most behavioral variance (4.2 unique sequences), followed by Gemini 3 Pro (3.0). All five models exhibit measurable behavioral inconsistency even on identical inputs.

Table 1. Performance across five models (200 tasks each, except GPT-4o and Gemini with 100–102). Seqs = unique action sequences; Var = step variance ratio.

Model	Correct	Seqs	Steps	Var
Claude Sonnet 4.5	**81.5%**	**2.3**	5.1	**22.2%**
GPT-5	79.5%	2.4	4.2	36.0%
Llama 3.1 70B	73.7%	4.2	4.7	64.9%
GPT-4o	73.3%	2.5	4.6	30.2%
Gemini 3 Pro	72.9%	3.0	5.2	36.3%

4.2. Consistency Is Associated with Correctness

Table 2 shows the central finding: consistent tasks ($\leq$2 unique sequences) achieve 80–87% accuracy while inconsistent tasks ($\geq$4 sequences) achieve 43–65%. The gap is significant for all five models (Mann-Whitney U, all $p < 0.001$; rank-biserial $r = -0.35$ to -0.58).

Table 2. Correctness (%) by consistency level. Gap in pp. r = rank-biserial; n = consistent/inconsistent counts. All $p < 0.001$.

Model	Cons. ($\leq$2)	Incons. ($\geq$4)	Gap	r	n
Gemini 3 Pro	86.5	44.7	41.8	−.58	65/34
GPT-4o	80.1	44.2	35.9	−.47	70/19
GPT-5	85.9	53.1	32.8	−.43	141/29
Claude Son. 4.5	86.4	60.3	26.1	−.39	155/32
Llama 70B	82.3	64.7	17.6	−.35	61/104

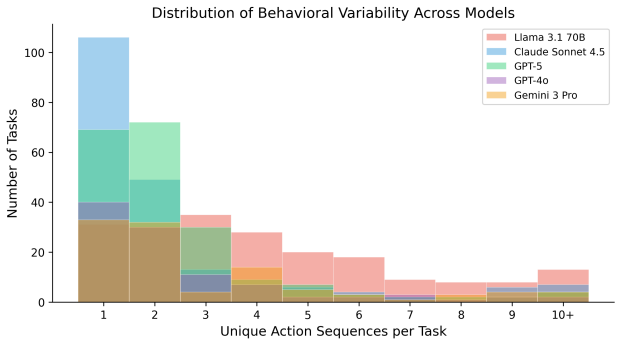

(a) Distribution of unique action sequences per task.

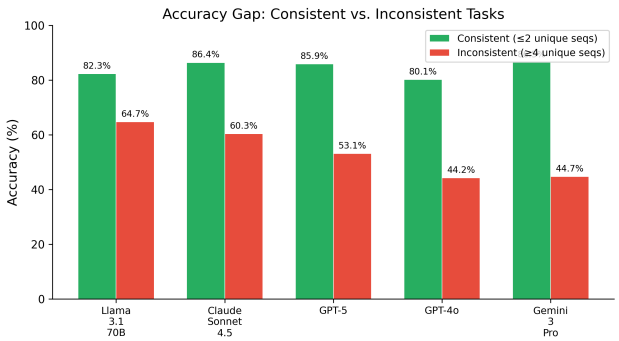

(b) Correctness for consistent ($\leq$2 seqs) vs. inconsistent ($\geq$4 seqs) tasks.

Figure 1. Behavioral consistency varies across models (left) and is associated with correctness (right).

Controlling for difficulty. Task difficulty confounds the consistency–correctness relationship. We compute a difficulty proxy (mean correctness across all models) and bin questions into Easy ($\geq$0.80, $n = 135$), Medium (0.40–0.79, $n = 31$), and Hard ($<$0.40, $n = 34$). The consistency gap persists within Easy and Medium strata (2–22pp) but reverses for Hard tasks, where inconsistent tasks marginally outperform consistent ones. This reversal is driven by *consistent-wrong* behavior: across models, 91% (53/58) of consistent Hard tasks commit to an incorrect answer, while 42% (27/65) of inconsistent Hard tasks recover at least one correct run through exploration (Appendix C). Partial correlations remain significant for three models: Gemini ($r = -0.50$, $p < 0.001$), Llama ($r = -0.18$, $p = 0.012$), and GPT-5 ($r = -0.17$, $p = 0.019$).

4.3. Consistency as a Runtime Failure Detector

We frame failure detection as binary classification: predicting whether the majority answer is incorrect using only consistency metrics. Table 3 reports the best feature per model.

Table 3. Failure detection via consistency metrics. H = answer entropy. Baseline precision equals the per-model error rate.

Model	Feat.	Prec.	Rec.	AUROC
Gemini 3 Pro	H>1.5	38.1%	47.1%	.689
GPT-5	H>1.5	33.3%	56.8%	.686
Llama 70B	H>1.5	35.3%	48.6%	.671
GPT-4o	Seq>4	66.7%	33.3%	.600
Claude Son. 4.5	H>0.5	17.0%	57.7%	.570

For the most variable models (Llama, GPT-5, Gemini), answer entropy achieves AUROC 0.67–0.69, substantially above the 0.50 random baseline, validating consistency as a lightweight runtime failure signal.

4.4. Majority Voting and Selective Prediction

Can consistency be *exploited*, not just observed? We simulate a majority-vote intervention using existing multi-run data: for budget k, subsample k runs per question, take the majority answer, and evaluate accuracy (500 bootstrap iterations).

Table 4. Majority-vote accuracy (%) and selective prediction. $k=1$ is single-run baseline. Sel. = unanimous $k=3$ (answer only when all 3 agree).

Model	$k=1$	$k=3$	Gain	Sel. Acc	Cov.
Llama 70B	73.6	75.6	+2.0	87.2	53.8
Claude Son. 4.5	81.5	82.2	+0.7	87.8	61.8
GPT-5	79.6	80.1	+0.5	88.1	55.2
GPT-4o	73.4	73.3	−0.0	72.3	11.7
Gemini 3 Pro	72.8	72.8	−0.0	75.9	64.8

Table 4 reveals that consistency works far better as a *filter* than an *aggregator*—a striking contrast with single-turn self-consistency (Wang et al., 2023), where majority voting improves CoT accuracy by 5–17pp. In our multi-step agentic setting, voting yields only +0–2pp because agentic errors are *systematic*: an early wrong action commits the agent to an incorrect trajectory, producing consistent-wrong answers that voting cannot correct. *Selective prediction*—answering only when all $k=3$ runs agree—is the stronger intervention: for Llama, Claude, and GPT-5, unanimous agreement achieves 87–88% accuracy (6–14pp above baselines) at 54–62% coverage. The asymmetry is explained by a 2×2 taxonomy of task outcomes: consistent-correct tasks (25–68% per model) are retained; inconsistent-wrong tasks (6–18%) are filtered; the ceiling is set by *consistent-wrong* tasks (5–13%), where the model confidently errs. Notably, 61% of Claude's errors are consistent-wrong vs. only 25% of Llama's, explaining why Llama benefits most from voting (full taxonomy in Appendix G).

4.5. Divergence Occurs Early; Path Length Matters

Pooled across five models, 20.0% [95% CI: 17.3–22.9%] of tasks diverge by step 2, but this rate varies dramatically: Llama diverges early in 52.5% [45.6–59.3%] of tasks vs. 7.5% [4.6–12.0%] for Claude and GPT-5. Among Llama's early-diverging tasks, accuracy is 71.7% vs. 85.8% for late-diverging tasks. The first search query largely determines the trajectory.

Path length correlates with both consistency and correctness: Spearman ρ between mean steps and unique sequences is 0.74 (Llama), 0.50 (Gemini), 0.48 (Claude), 0.47 (GPT-5), 0.43 (GPT-4o); pooled $\rho = 0.48$, all $p < 0.001$. Longer paths indicate backtracking and uncertainty; each additional step is an opportunity to diverge and err.

4.6. Temperature Ablation

Figure 2 shows results across four models (Gemini temperature ablation partially complete; see Appendix A). All models show monotonically increasing unique sequences with temperature. Claude achieves near-perfect consistency at $T = 0.0$ (1.0 sequences) while Llama retains 2.3, confirming that model architecture matters beyond sampling noise. The accuracy–temperature tradeoff is model-dependent: GPT-4o gains +7pp at $T = 0.0$, while GPT-5 and Llama show minimal accuracy change. The finding that $T = 0.0$ does not eliminate divergence for Llama or GPT-5 indicates that agentic inconsistency arises from factors beyond token-level sampling. Full table in Appendix B.

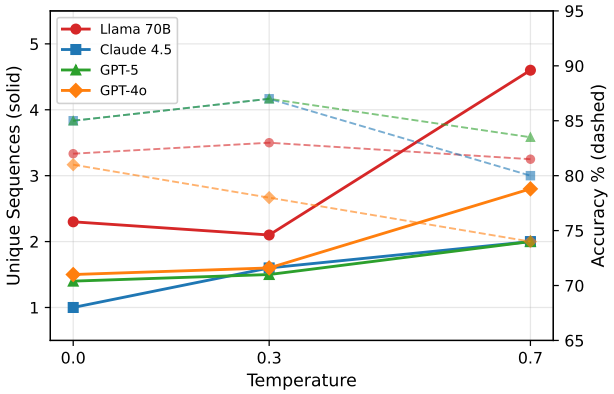

Figure 2. Temperature ablation (20 questions, 5 runs). Solid: unique sequences (left axis); dashed: accuracy (right axis).

4.7. Cross-Benchmark Pilot: SWE-bench

To test generalizability beyond HotpotQA's minimal action space, we conduct a pilot study on SWE-bench Verified (Jimenez et al., 2023a; Chow et al., 2024): 10 astropy tasks, 5 runs per model-task pair (150 runs), using Claude, GPT-5, and Llama with an identical bash-only scaffold.

Table 5. SWE-bench pilot (10 tasks × 5 runs). CV = std/mean of step counts. The HotpotQA consistency hierarchy is preserved.

	Claude	**GPT-5**	**Llama**
CV	**15.2%**	32.2%	47.0%
Accuracy	**58%**	32%	4%
Unique Seqs	100%	100%	100%

The consistency ranking is identical to HotpotQA (Claude > GPT-5 > Llama), consistent with consistency being a stable model-level property. All 150 runs produced unique action sequences, yet CV still discriminates between models. The accuracy gap widens from 8pp (HotpotQA) to 54pp (SWE-bench), consistent with the hypothesis that complexity amplifies consistency effects. Given the small sample (10 tasks, single repository), these results are preliminary; full cross-benchmark comparison in Appendix D.

4.8. Ranking Instability

We perform 10,000 bootstrap iterations across five models on the 100 shared questions, sampling one run per question per model. **70.1% [69.2–71.0%] of single-run evaluations produce a ranking differing from the multi-run ground truth** (Claude 81.0% > GPT-5 80.2% > Llama 76.6% > GPT-4o 73.3% = Gemini 73.3%). The instability stems from narrow accuracy gaps and within-model variance; the tied bottom-two models particularly amplify reordering. Agent benchmarks should report multi-run statistics.

5. Discussion

Beyond τ-bench. While τ-bench (Yao et al., 2024) established that agents are inconsistent, we provide complementary insights: *where* variance originates (step 2), its association with correctness ($r = -0.35$ to -0.58), and its practical utility for failure detection (AUROC 0.67–0.69). These findings suggest actionable interventions beyond measuring pass^k.

Practical implications. Consistency monitoring is feasible and actionable. Our majority-vote experiment (Section 4.4) shows that even $k = 3$ runs suffice: selective prediction on unanimous agreement achieves 87–88% accuracy for three models, a 6–14pp gain over single-run baselines at 54–62% coverage. Crucially, consistency functions as a *filter* (abstain on disagreement) rather than an *aggregator* (vote to pick the best): voting gains are minimal (+0–2pp), but filtering gains are substantial (+4–13pp). The ceiling on filtering is set by consistent-wrong tasks (5–13% per model), where the agent confidently commits to an incorrect answer. Answer entropy provides a complementary lightweight signal (AUROC 0.67–0.69). The step 2 bottleneck suggests that improving first-query formulation could

reduce downstream variance.

Consistency as implicit calibration. Traditional calibration asks whether a model's stated confidence matches its accuracy (Kadavath et al., 2022). Behavioral consistency provides an alternative, *behavioral* notion of calibration: agents that produce diverse trajectories across runs are implicitly signaling uncertainty about the task, much as deep-ensemble disagreement signals epistemic uncertainty (Lakshminarayanan et al., 2017). Unlike verbalized confidence—which requires prompting strategies and is often poorly calibrated (Xiong et al., 2024; Tian et al., 2023)—behavioral consistency emerges naturally from repeated execution and requires no model modification. Our selective-prediction results (Table 4) confirm that this implicit signal is actionable: unanimous agreement among $k = 3$ runs selects for high-accuracy subsets. This positions multi-run consistency as a practical calibration mechanism for deployed agents.

Limitations and future work. Our analysis primarily uses HotpotQA with a 3-tool action space; the SWE-bench pilot provides preliminary evidence of generalizability. The temperature ablation uses a smaller sample (20 questions). We cannot establish causality from observational data. Gemini 3 Pro required a `top_p` workaround to achieve stochastic sampling (Appendix A), and its Q100–199 data collection is ongoing; current Gemini results use 102 questions. We plan extensions to multimodal agents and analytical reasoning (FinQA, GAIA).

6. Conclusion

Behavioral consistency in LLM agents is measurable, associated with correctness, and practically exploitable. Across 10,000 runs on five models, consistent agents achieve 80–87% accuracy vs. 43–65% for inconsistent ones, with effect sizes that survive difficulty controls. Consistency metrics detect failures with AUROC 0.67–0.69, and selective prediction on unanimous $k = 3$ agreement yields 87–88% accuracy at 54–62% coverage—demonstrating that behavioral consistency functions as an actionable confidence signal. Single-run evaluations produce incorrect rankings 70.1% of the time. These findings argue for multi-run evaluation, consistency-based selective prediction, and runtime consistency monitoring as standard practice.

Code and data: `https://github.com/amanmehta-maniac/agent-consistency`.

Impact Statement

Our finding that single-run evaluations produce incorrect model rankings 70% of the time has direct implications for practitioners who rely on benchmark scores to select agent models. For high-stakes deployments, monitoring behavioral consistency across parallel executions could serve as a lightweight failure detection mechanism with AUROC 0.67–0.69. We provide diagnostic tools for understanding existing limitations rather than new capabilities, and do not foresee negative societal consequences from this work.

References

Chow, N., Jimenez, C. E., Yang, J., and Narasimhan, K. Swe-bench verified: A human-validated subset for reliable agent evaluation. *OpenAI Blog*, 2024. Available at: `https://openai.com/index/introducing-swe-bench-verified/`.

Geifman, Y. and El-Yaniv, R. Selective classification for deep neural networks. *NeurIPS*, 2017.

Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., and Narasimhan, K. Swe-bench: Can language models resolve real-world github issues? *arXiv preprint arXiv:2310.06770*, 2023a.

Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., and Narasimhan, K. Swe-bench: Can language models resolve real-world github issues? *arXiv preprint arXiv:2310.06770*, 2023b.

Kadavath, S., Conerly, T., Askell, A., Henighan, T., Drain, D., Perez, E., Schiefer, N., Hatfield-Dodds, Z., DasSarma, N., Tran-Johnson, E., et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.

Kamath, A., Jia, R., and Liang, P. Selective question answering under domain shift. *ACL*, 2020.

Kuhn, L., Gal, Y., and Farquhar, S. Semantic uncertainty: Linguistic invariances for uncertainty estimation in natural language generation. *arXiv preprint arXiv:2302.09664*, 2023.

Lakshminarayanan, B., Pritzel, A., and Blundell, C. Simple and scalable predictive uncertainty estimation using deep ensembles. *NeurIPS*, 2017.

Schick, T., Dwivedi-Yu, J., Dessì, R., Raileanu, R., Lomeli, M., Zettlemoyer, L., Cancedda, N., and Scialom, T. Toolformer: Language models can teach themselves to use tools. *arXiv preprint arXiv:2302.04761*, 2023.

Stroebl, P. et al. Are language model benchmarks overestimating agent capabilities? *arXiv preprint*, 2024.

Tian, K., Mitchell, E., Yao, H., Manning, C. D., and Finn, C. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. *EMNLP*, 2023.

Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., Chowdhery, A., and Zhou, D. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171*, 2023.

Wu, Q., Bansal, G., Zhang, J., Wu, Y., Li, B., Zhu, E., Jiang, L., Zhang, X., Zhang, S., Liu, J., et al. Autogen: Enabling next-gen llm applications via multi-agent conversation. *arXiv preprint arXiv:2308.08155*, 2023.

Xiong, M., Hu, Z., Lu, X., Li, Y., Fu, J., He, J., and Hooi, B. Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs. *ICLR*, 2024.

Yang, Z., Qi, P., Zhang, S., Bengio, Y., Cohen, W. W., Salakhutdinov, R., and Manning, C. D. Hotpotqa: A dataset for diverse, explainable multi-hop question answering. *arXiv preprint arXiv:1809.09600*, 2018.

Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., and Cao, Y. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*, 2022.

Yao, S., Narasimhan, K., and Cao, Y. τ-bench: A benchmark for tool-agent-user interaction in real-world domains. *arXiv preprint arXiv:2406.12045*, 2024.

Zhou, S., Xu, F. F., Zhu, H., Zhou, X., Lo, R., Sridhar, A., Cheng, X., Bisk, Y., Fried, D., Alon, U., et al. Webarena: A realistic web environment for building autonomous agents. *arXiv preprint arXiv:2307.13854*, 2023.

A. Gemini 3 Pro: Sampling Configuration

Gemini 3 Pro on the Snowflake Cortex endpoint initially exhibited near-deterministic outputs at $T = 0.7$, producing identical action sequences in >99% of runs. We determined that the endpoint does not fully respect the `temperature` parameter without explicitly setting `top_p` < 1.0. After adding `top_p = 0.95`, outputs became properly stochastic (3.0 unique sequences on average, comparable to GPT-4o and GPT-5). Main results use the corrected stochastic data. Temperature ablation for Gemini is partially complete ($T = 0.3$: Acc 76.0%, 2.0 unique sequences; $T = 0.7$: Acc 76.5%, 2.5 unique sequences); $T = 0.0$ data collection is ongoing.

B. Temperature Ablation: Full Results

Table 6. Full temperature ablation (20 questions, 5 runs each). $T = 0.7$ uses same questions from the main experiment.

Model	Temp	Accuracy	Unique Seqs
Llama 3.1 70B	0.0	82.0%	2.3
	0.3	83.0%	2.1
	0.7	81.5%	4.6
Claude Sonnet 4.5	0.0	85.0%	1.0
	0.3	87.0%	1.6
	0.7	80.0%	2.0
GPT-5	0.0	85.0%	1.4
	0.3	87.0%	1.5
	0.7	83.5%	2.0
GPT-4o	0.0	81.0%	1.5
	0.3	78.0%	1.6
	0.7	74.0%	2.8

C. Difficulty Stratification: Full Results

Task difficulty is computed as mean correctness across all runs from all models. Strata: Easy (≥0.80), Medium (0.40–0.79), Hard (<0.40). Of 200 questions: 135 Easy, 31 Medium, 34 Hard.

Table 7. Consistency–correctness gap (pp) within difficulty strata.

Stratum	Llama	GPT-4o	GPT-5	Claude	Gemini
Easy (135)	9.2pp	1.0pp	5.9pp	1.5pp	17.8pp
Medium (31)	−0.5pp	3.1pp	15.8pp	3.3pp	22.3pp
Hard (34)	−13.3pp	−7.1pp	1.2pp	−5.3pp	−5.5pp

Partial correlations between unique sequences and accuracy, controlling for difficulty (continuous): Gemini $r = -0.496$ ($p < 0.001$), Llama $r = -0.178$ ($p = 0.012$), GPT-5 $r = -0.166$ ($p = 0.019$), GPT-4o $r = -0.075$ ($p = 0.458$), Claude $r = -0.035$ ($p = 0.620$). The association remains significant for three of five models.

Consistent-Wrong analysis. The Hard stratum reversal (negative gaps above) is explained by *consistent-wrong* behavior: on tasks where the model lacks knowledge, consistency reflects confident commitment to an incorrect answer rather than reliable competence. We classify each model×Hard-task pair into four categories based on unique action sequences and correctness:

- **Consistent-Correct** (≤2 seqs, majority correct): 5/137 (3.6%)
- **Consistent-Wrong** (≤2 seqs, majority wrong): 53/137 (38.7%)
- **Inconsistent-Lucky** (≥4 seqs, ≥1 correct run): 27/137 (19.7%)
- **Inconsistent-Wrong** (≥4 seqs, 0 correct): 38/137 (27.7%)

Among consistent Hard tasks, 91.4% (53/58) are Consistent-Wrong—the model locks into one incorrect reasoning path. Conversely, 41.5% (27/65) of inconsistent Hard tasks achieve at least one correct run (mean 3.1/10 correct runs), confirming that behavioral variability enables exploration that occasionally recovers the right answer. This also partly reflects correctness metric sensitivity: several Consistent-Wrong tasks have semantically correct model answers (e.g.,

"Pasek and Paul" for gold "Pasek & Paul", "Kohlberg Kravis Roberts" for gold "KKR & Co") that fail fuzzy matching. Per-model consistent-wrong rates: GPT-5 94% (16/17), Llama 100% (7/7), Claude 79% (15/19), Gemini 100% (7/7), GPT-4o 100% (8/8).

D. Cross-Benchmark Comparison

Table 8. HotpotQA vs. SWE-bench. Consistency ranking preserved; accuracy gaps widen with task complexity.

	HotpotQA	SWE-bench
Action space	3 tools	Unrestricted bash
Avg trajectory	3–8 steps	10–46 steps
Best model CV	8.2% (Claude)	15.2% (Claude)
Worst model CV	22.1% (Llama)	47.0% (Llama)
Accuracy range	73–82%	4–58%
Mean unique seqs	2.3–4.2	all unique
Ranking preserved	✓ (Claude > GPT > Llama)	

E. Question Type Analysis

On the first 100 questions with type annotations, we compare bridge questions (multi-hop, $n = 79$) with comparison questions (yes/no, $n = 21$) using Llama 3.1 70B. Comparison questions show *higher* correctness (80.0% vs. 75.7%) but *lower* answer consistency (62.4% vs. 76.6%) and lower step variance (41% vs. 63%), highlighting that answer consistency and explanation consistency are distinct dimensions.

F. Metric Sensitivity

We verify robustness to the choice of correctness metric. Under exact match (EM), overall accuracy decreases (e.g., Claude: 43.6% EM vs. 81.5% fuzzy) but model rankings and consistency gaps are preserved. Under token F1 $>$ 0.5, results closely track fuzzy match. The consistency–correctness association holds across all three metrics for all five models.

G. Majority Voting and Selective Prediction: Full Results

Task-outcome taxonomy. Table 9 classifies each task by consistency ($\leq$2 vs. $\geq$4 unique sequences) and correctness (majority answer matches gold). The taxonomy explains why filtering outperforms voting: consistent-wrong tasks (5–13%) set the ceiling on filtering, while inconsistent-correct tasks (8–38%)—where voting could theoretically consolidate the right answer—are rare except for Llama.

Table 10 reports majority-vote accuracy across all budget levels. Gains are largest for Llama (+4.4pp at $k = 10$) and Claude (+2.0pp), while GPT-4o and Gemini show no improvement—consistent with the hypothesis that voting helps only when the model explores diverse-but-sometimes-correct paths.

Table 9. Task-outcome taxonomy (% of tasks). Consistent: $\leq$2 unique seqs; Inconsistent: $\geq$4.

Model	Cons-Right	Cons-Wrong	Inc-Right	Inc-Wrong
Claude Son. 4.5	67.5	10.0	10.5	5.5
GPT-5	61.5	9.0	7.5	7.0
GPT-4o	57.0	13.0	7.0	12.0
Gemini 3 Pro	52.1	9.2	15.1	17.6
Llama 70B	25.0	5.5	38.0	14.0

Table 10. Majority-vote accuracy (%) by budget k (500 bootstrap iterations). Gain is relative to $k = 1$.

Model	$k=1$	$k=3$	$k=5$	$k=7$	$k=10$
Llama 70B	73.6	75.6+2.0	76.9+3.3	77.2+3.6	78.0+4.4
Claude S. 4.5	81.5	82.2+0.7	82.6+1.0	82.9+1.4	83.5+2.0
GPT-5	79.6	80.1+0.5	80.2+0.6	80.3+0.7	80.5+0.9
GPT-4o	73.4	73.3−0.0	73.2−0.2	72.8−0.6	73.0−0.4
Gemini 3 Pro	72.8	72.8−0.0	72.4−0.4	72.2−0.6	71.6−1.2

Table 11 reports selective prediction at multiple agreement thresholds. Stricter thresholds yield higher accuracy but lower coverage. The accuracy–coverage tradeoff is favorable for Llama, Claude, and GPT-5; GPT-4o's extremely low coverage under unanimity reflects its high answer variability.

Table 11. Selective prediction: accuracy (%) and coverage (%) at varying agreement thresholds.

Model	Threshold	Acc.	Cov.
Llama 70B	Unanimous 3/3	87.2	53.8
	Majority 2/3	82.2	82.5
	Unanimous 5/5	88.7	44.2
	4/5	86.7	63.8
	Majority 3/5	83.6	80.3
Claude Son. 4.5	Unanimous 3/3	87.8	61.8
	Majority 2/3	84.4	84.4
	Unanimous 5/5	88.4	55.4
	4/5	87.8	67.3
	Majority 3/5	85.2	82.5
GPT-5	Unanimous 3/3	88.1	55.2
	Majority 2/3	83.2	78.8
	Unanimous 5/5	90.2	48.1
	4/5	86.9	62.3
	Majority 3/5	84.1	75.3
GPT-4o	Unanimous 3/3	72.3	11.7
	Majority 2/3	78.0	42.4
	Unanimous 5/5	64.8	5.7
	4/5	73.3	16.1
	Majority 3/5	78.7	35.0
Gemini 3 Pro	Unanimous 3/3	75.9	64.8
	Majority 2/3	73.3	84.8
	Unanimous 5/5	77.1	58.6
	4/5	75.3	71.3
	Majority 3/5	73.5	82.1

Table 12 shows how the $k = 3$ majority-vote gain varies by difficulty stratum. Gains concentrate on Easy and Medium tasks; Hard tasks show no improvement, consistent with the

finding that consistency on hard tasks reflects consistent-wrong behavior (Appendix C).

Why self-consistency works less in agentic settings. Wang et al. (2023) report 5–17pp gains from majority voting on single-turn CoT tasks (arithmetic, commonsense). Our agentic setting shows only 0–2pp gains. The difference is structural: in single-turn CoT, errors are primarily in the final reasoning step and the answer space is small, so diverse chains often converge on the correct answer. In multi-step agentic settings, an early wrong action (e.g., searching for the wrong entity at step 2) commits the agent to an incorrect trajectory for all subsequent steps, producing systematic errors that voting cannot correct. This makes consistency more valuable as a *diagnostic* signal (is the agent confident?) than as an *ensembling* strategy (pick the best answer).

Table 12. Per-stratum $k=3$ majority-vote gain (pp) over $k=1$.

Stratum	Llama	Claude	GPT-5	GPT-4o	Gemini
Easy	+2.4	+0.3	+0.7	+0.3	+0.7
Medium	+2.4	+1.8	+1.8	−0.8	−1.2
Hard	−0.0	+0.9	−1.4	+0.1	−0.5